

IIT Kharagpur



IIT Madras



IIT Goa



IIT PALAKKAD

IIT Palakkad

Applied Accelerated AI

Scale training on multiple GPUs with PyTorch

Dr. Satyajit Das

Assistant Professor

Data Science

Computer Science and Engineering

IIT Palakkad



National
Supercomputing
Mission



Centre for
Development of
Advanced Computing



Scale training to multiple GPUs


- **Data Parallel** – Data distributed across devices
- **Model Parallel** – Model distributed across devices



Scale training to multiple GPUs

- Single Machine Data Parallel
- Single Machine Model Parallel
- Distributed Data Parallel
- Distributed Data Parallel with Model Parallel
- Distributed Model Parallel

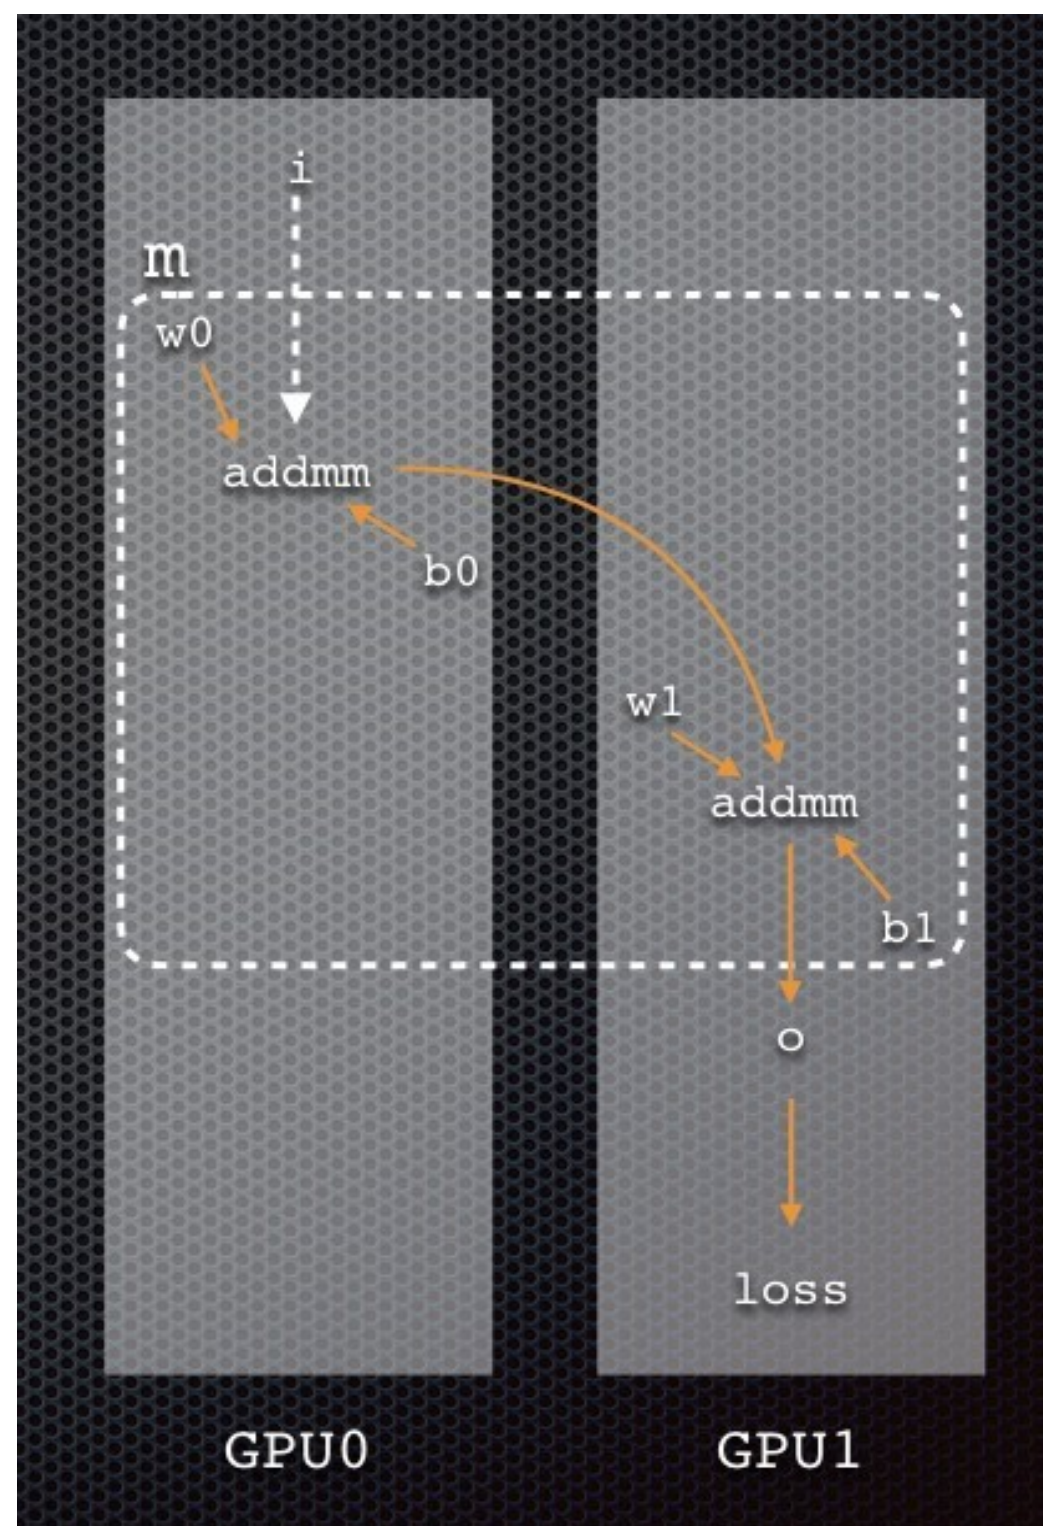
Single Machine Data Parallel



```
model = Net().to("cuda:0")
model = torch.nn.DataParallel(model)

# training loop ...
```


Single Machine Model Parallel



```
class Net(torch.nn.Module):
```

```
    def __init__(self, *gpus):
        super(Net).__init__(self)
```

```
        self.gpu0 = torch.device(gpus[0])
```

```
        self.gpu1 = torch.device(gpus[1])
```

```
        self.sub_net1 = torch.nn.Linear(10, 10).to(self.gpu0)
```

```
        self.sub_net2 = torch.nn.Linear(10, 5).to(self.gpu1)
```

```
    def forward(self, x):
```

```
        y = self.sub_net1(x.to(self.gpu0))
```

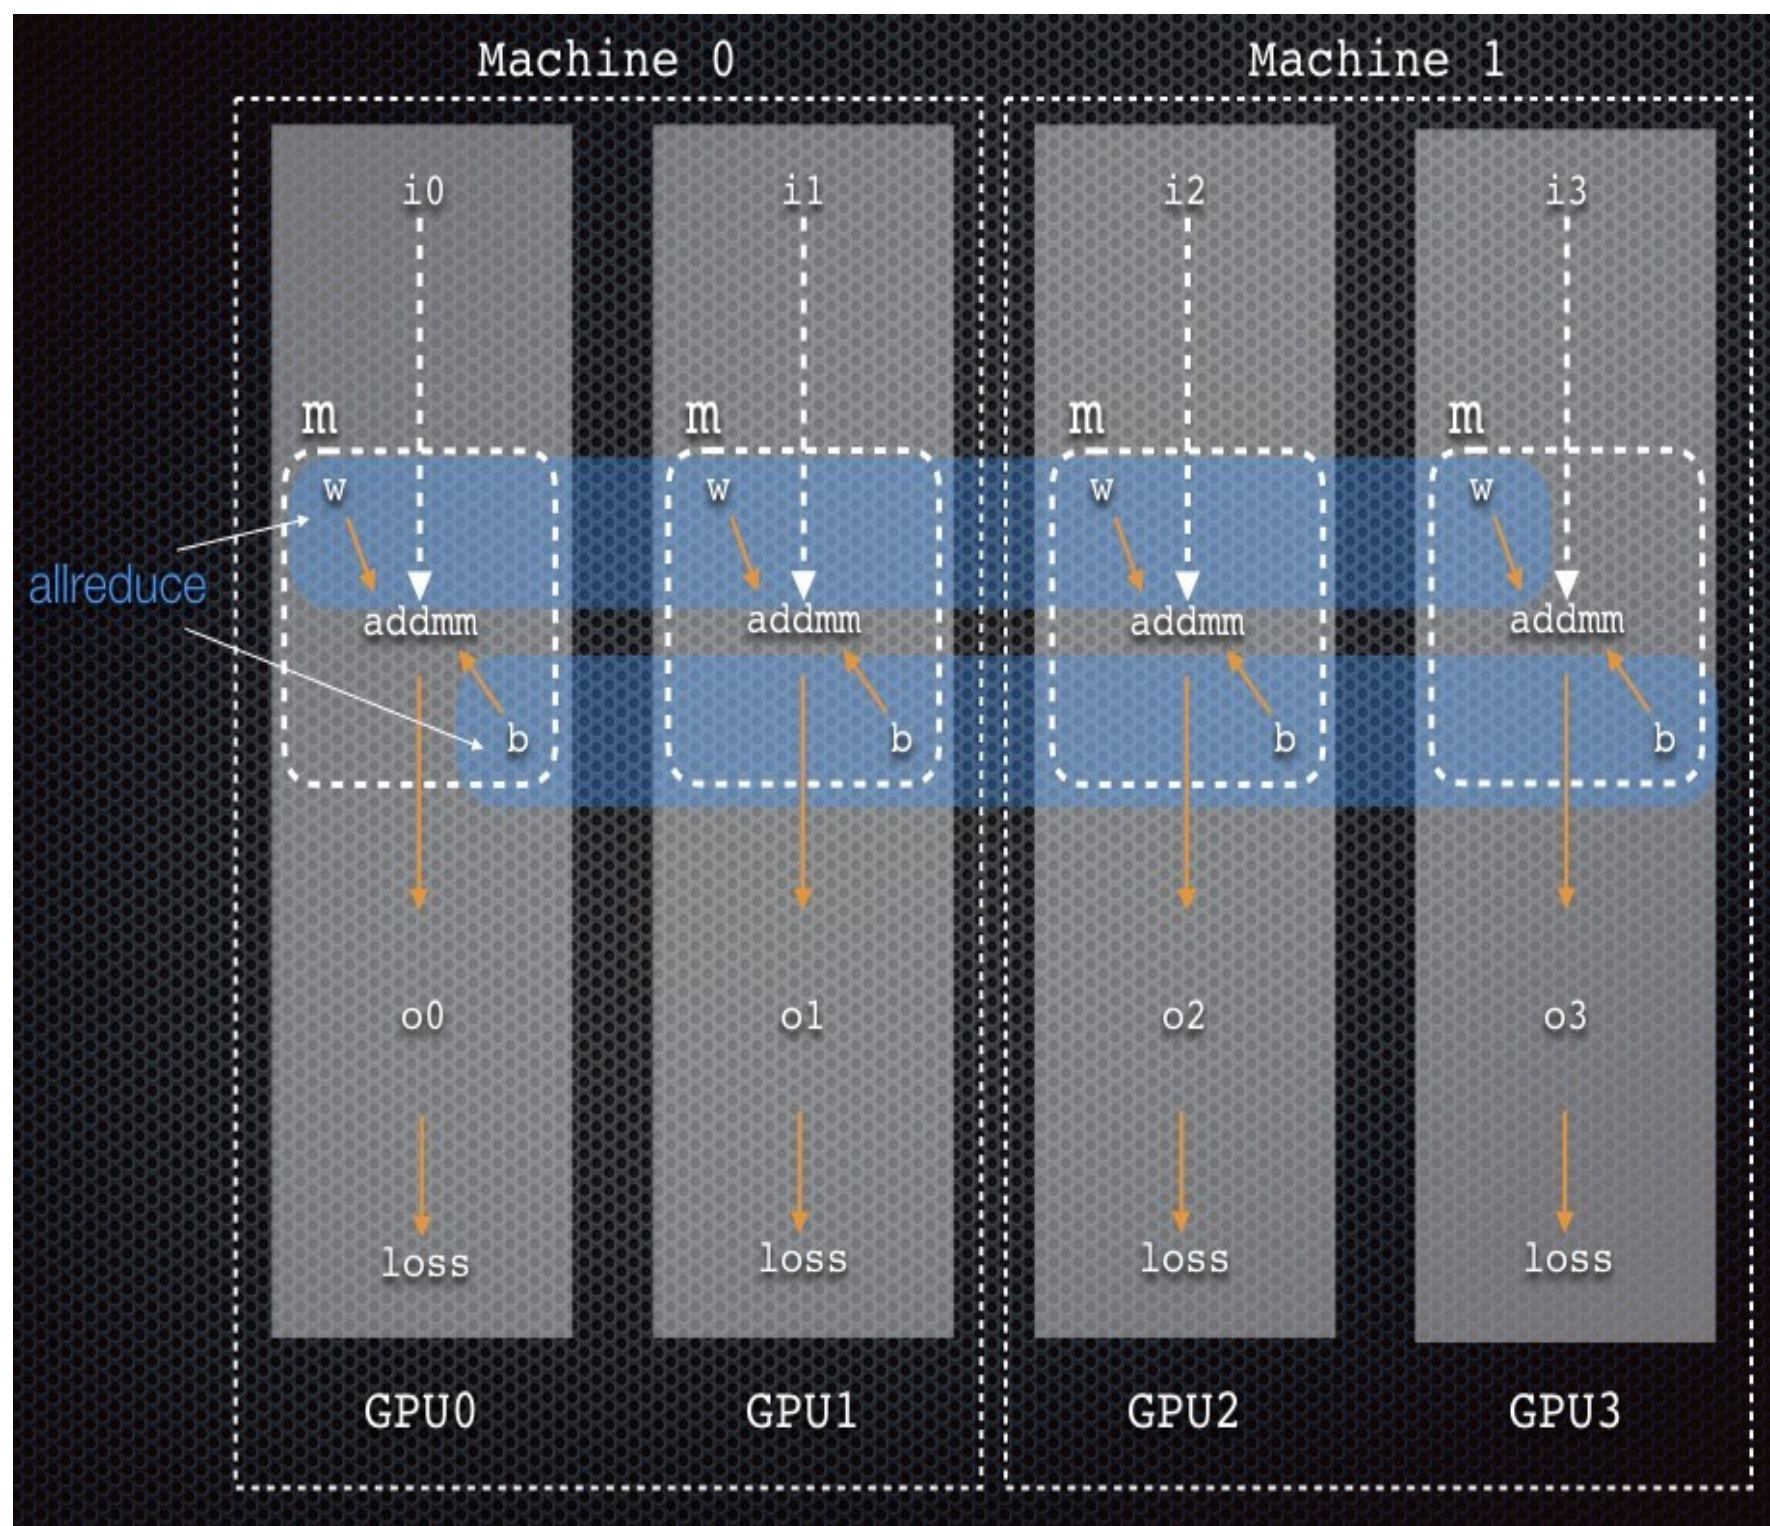
```
        z = self.sub_net2(y.to(self.gpu1)) # blocking
```

```
        return z
```

```
model = Net("cuda:0", "cuda:1")
```

```
# training loop...
```


Distributed Data Parallel



```
def one_machine(machine_rank, world_size,
                backend): torch.distributed.init_process_group(
                    backend, rank=machine_rank, world_size=world_size
                )

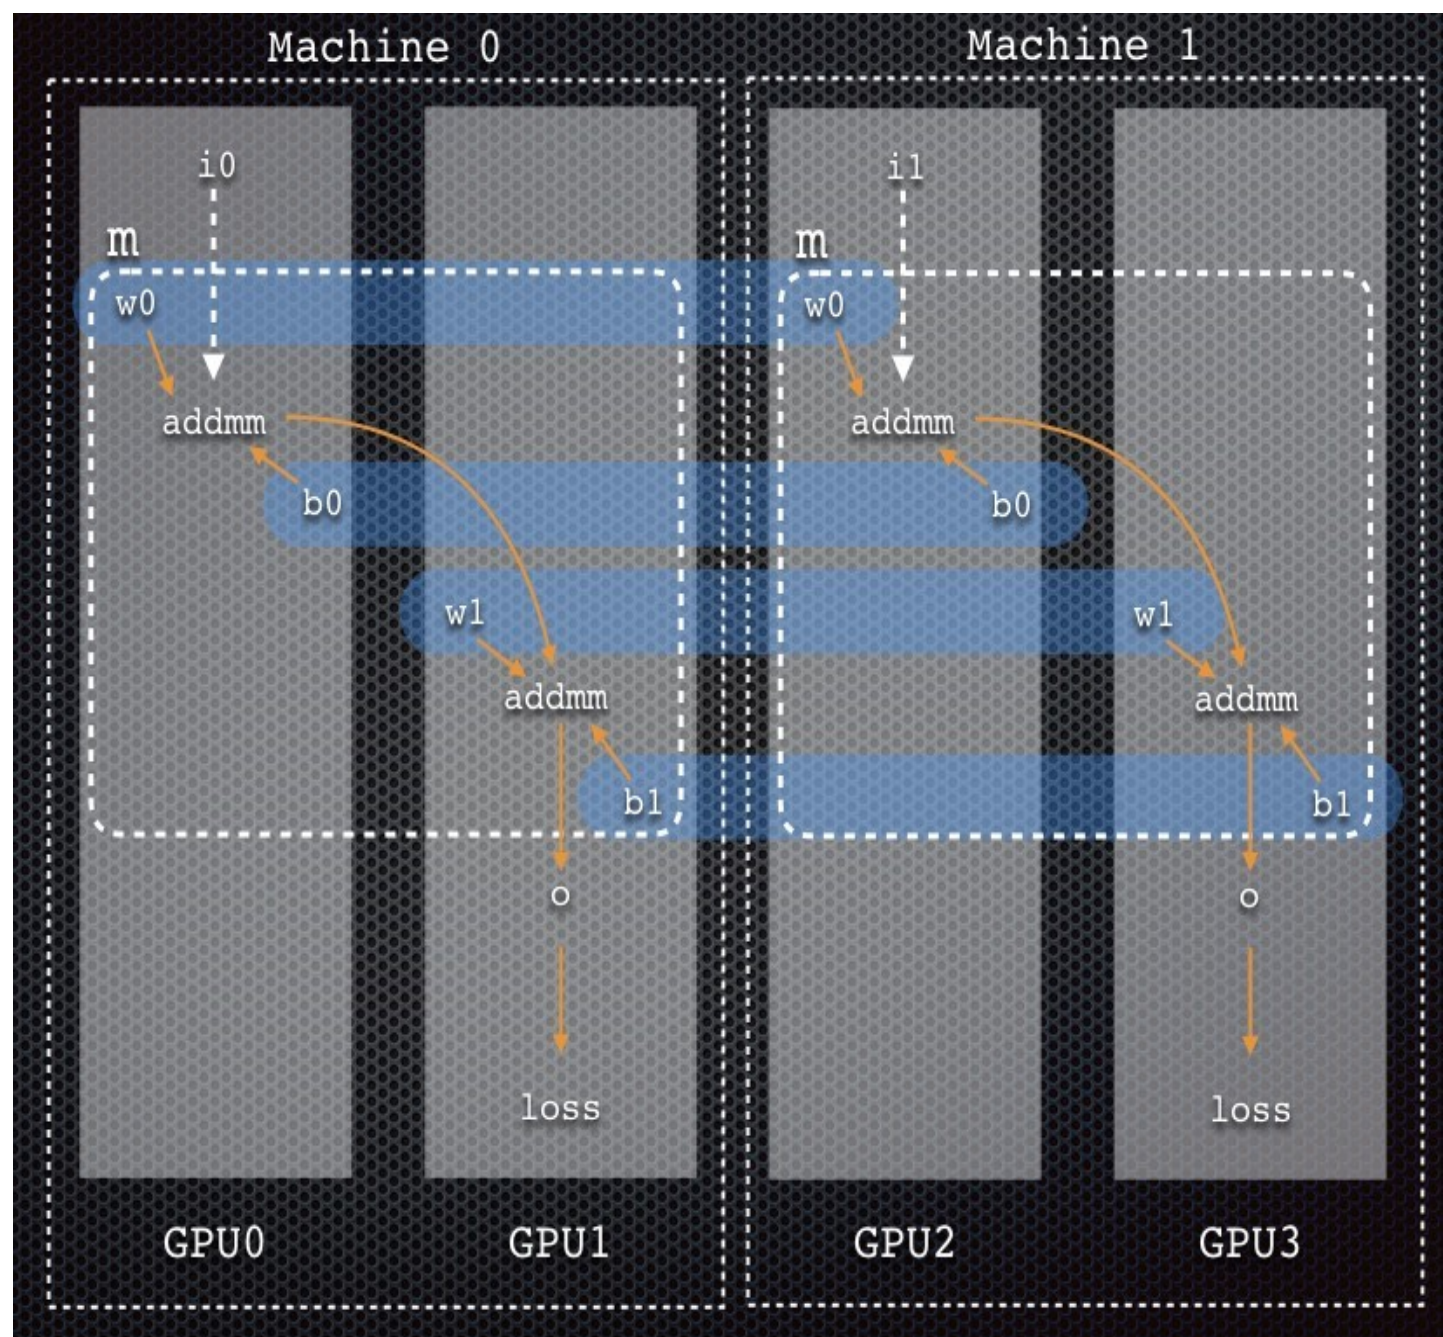
gpus = {
    0: [0, 1],
    1: [2, 3],
}[machine_rank] # or one gpu per process to avoid GIL

model = Net().to(gpus[0]) # default to first gpu on
machine
model = torch.nn.parallel.DDP(model, device_ids=gpus)

# training loop...

for machine_rank in range(world_size):
    torch.multiprocessing.spawn(
        one_machine, args=(world_size, backend),
        nprocs=world_size, join=True # blocking
    )
```


Distributed Data Parallel with Model Parallel



```
def one_machine(machine_rank, world_size, backend):
    torch.distributed.init_process_group(
        backend, rank=machine_rank, world_size=world_size
    )
```

```
gpus = {0: [0, 1],
        1: [2, 3],}[machine_rank]
```

```
model = Net(gpus)
model = torch.nn.parallel.DDP(model)
```

```
# training loop...
```

```
for machine_rank in range(world_size):
    torch.multiprocessing.spawn(
        one_machine, args=(world_size, backend),
        nprocs=world_size, join=True
    )
```

Thank You