

IIT Kharagpur



IIT Madras



IIT Goa



IIT PALAKKAD

IIT Palakkad

APPLIED ACCELERATED ARTIFICIAL INTELLIGENCE

TensorFlow

Dr. Satyajit Das

Assistant Professor

Data Science

Computer Science and Engineering

IIT Palakkad



National
Supercomputing
Mission


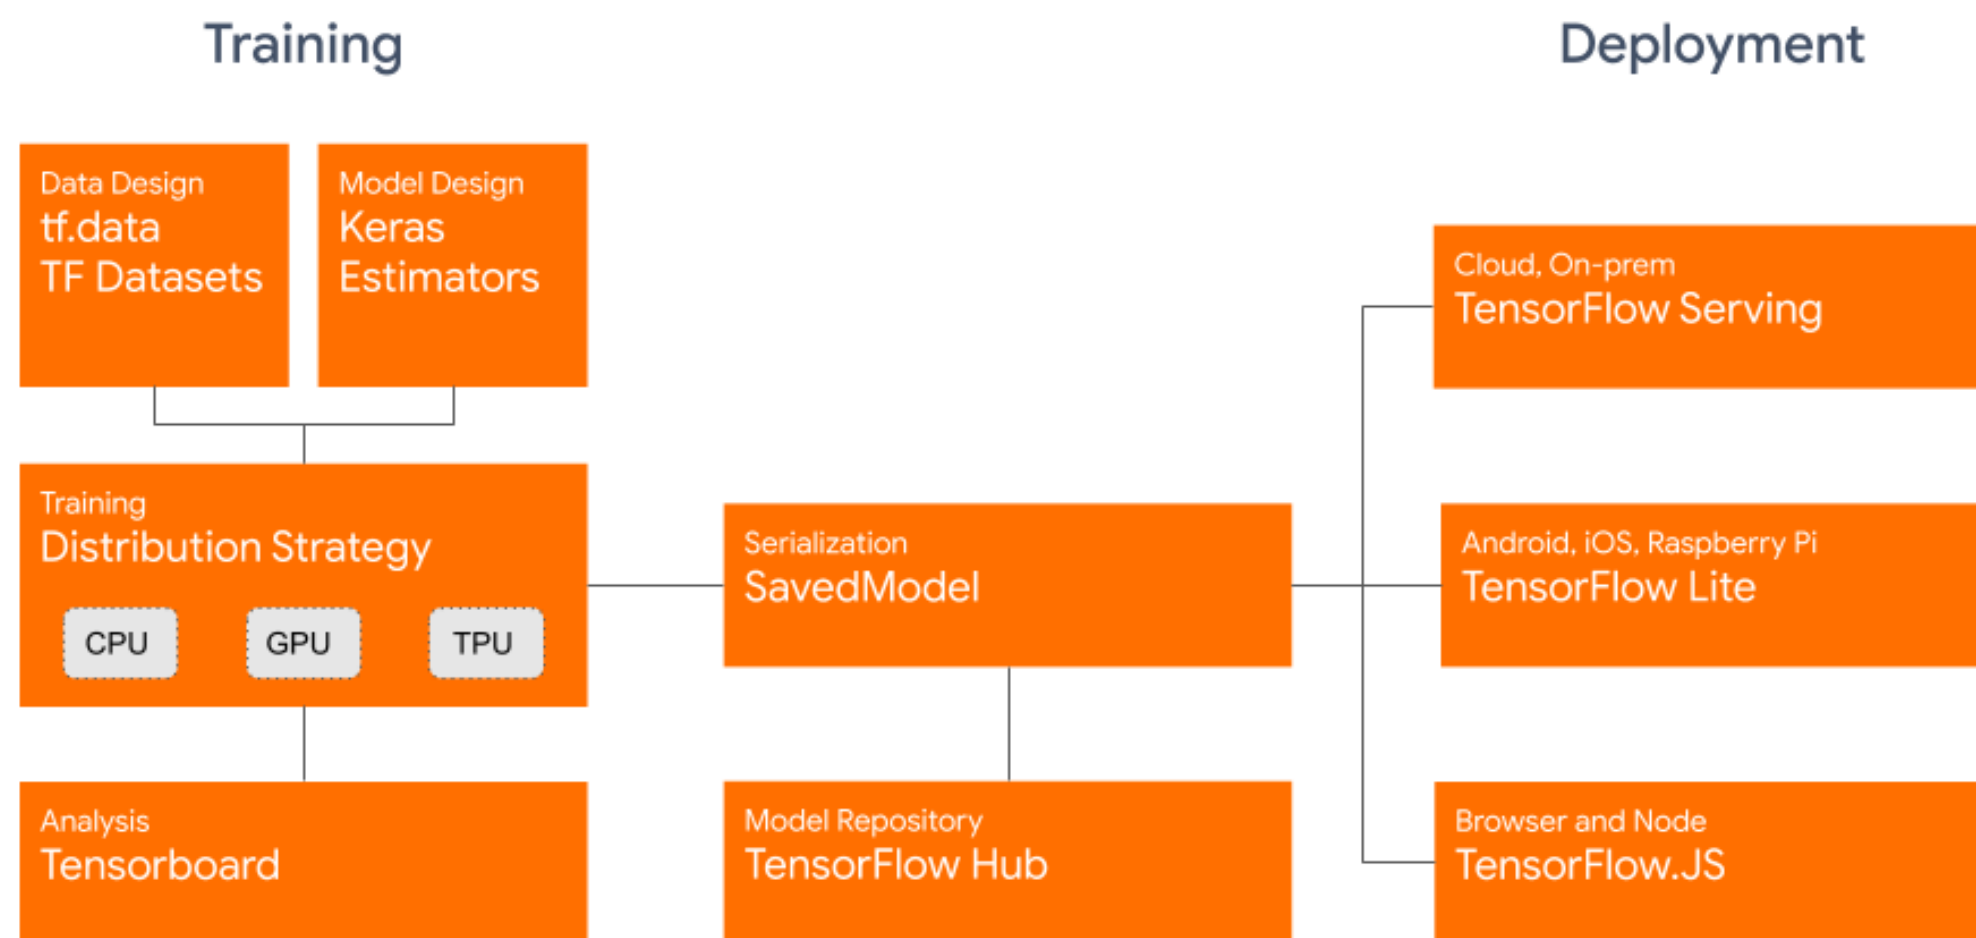


Centre for
Development of
Advanced Computing



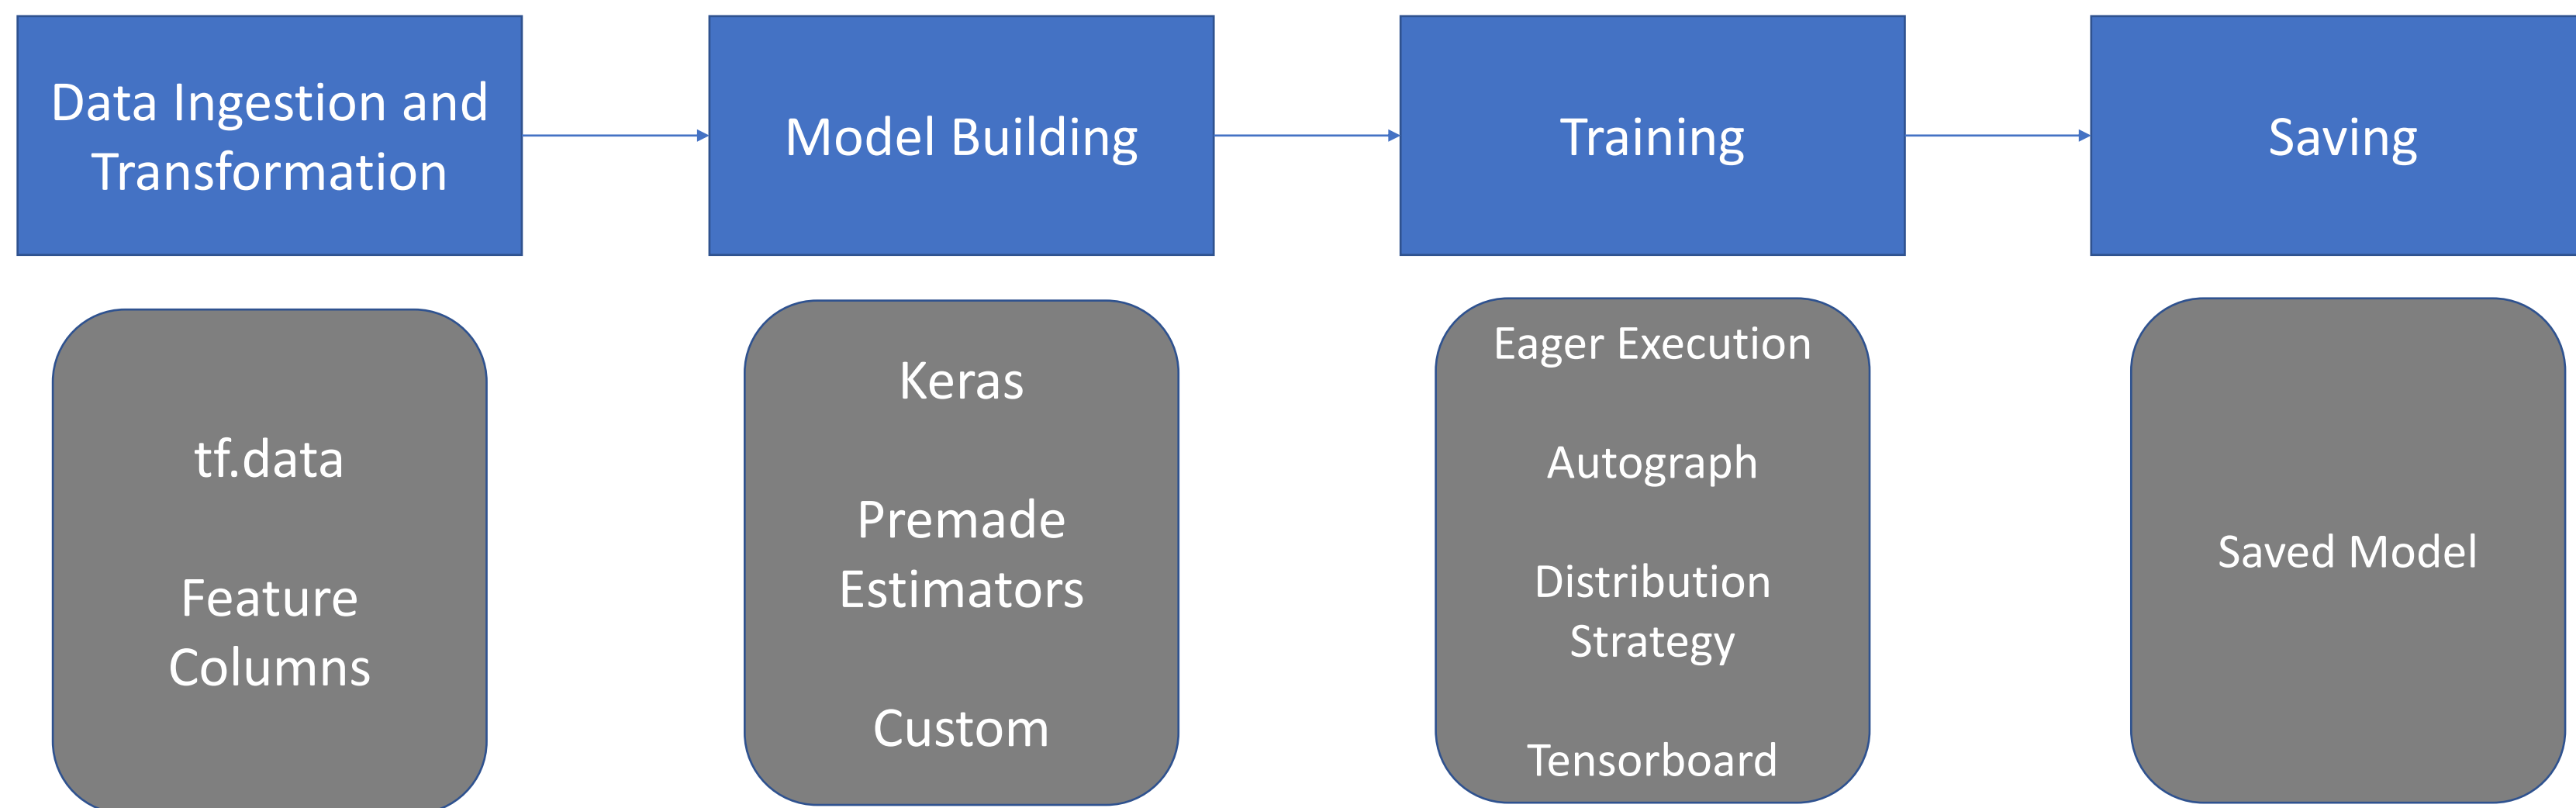
TensorFlow 1.x/2.x

- An open source Deep Learning library
 - >1,800 contributors worldwide
 - Apache 2.0 license
 - Released by Google in 2015
- TensorFlow 2.0
 - Easier to learn and use
 - For beginners and experts
 - Available today





The workflow



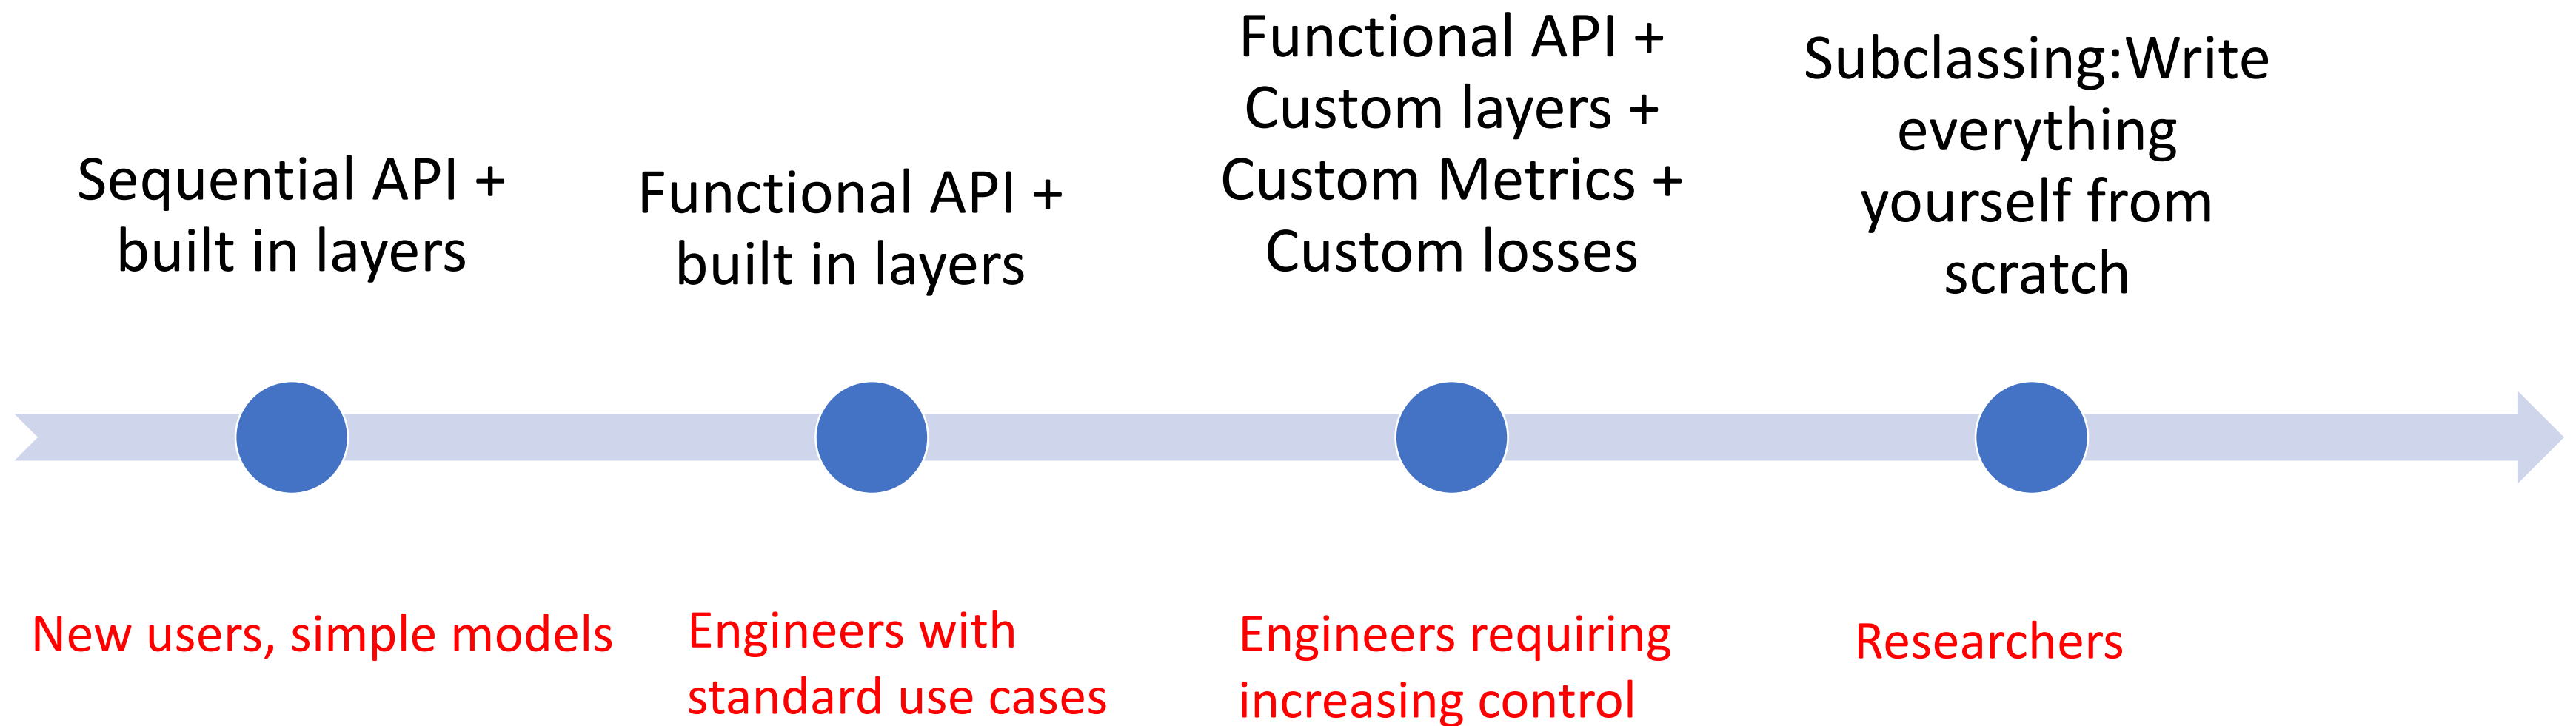


What's new in TensorFlow 2.x

- Easy model building with **Keras and eager execution** (activated by default in TF2.0).
- Robust model **deployment** in production on any platform.
- Powerful experimentation for research.
- Simplifying the API by **cleaning up deprecated APIs** and reducing duplication (relevant in case you have code developed in TensorFlow 1.X and you need to convert it)
- Load your data using **tf.data**. Training data is read using input pipelines which are created using tf.data.
- Build, train and validate your model with **tf.keras**, or use **Premade Estimators**.
- **TensorFlow Hub**.
- **Run and debug with eager execution**, then use tf.function for the benefits of graphs.
- Use Distribution Strategies for distributed training.
- hardware accelerators like CPUs, GPUs, and TPUs; you can enable training workloads to be distributed to single-node/multi-accelerator as well as multi-node/multi-accelerator configurations, including TPU Pods.
- Export to **SavedModel**. TensorFlow will standardize on SavedModel as an interchange format for TensorFlow Serving, TensorFlow Lite, TensorFlow.js, TensorFlow Hub, and more.
- Tensorflow Datasets



Model Building





Symbolic vs Imperative APIs

- Symbolic (Keras Sequential)
 - Your model is a graph of layers
 - Any graph you compile will run
 - TensorFlow helps you debug by catching errors at compile time

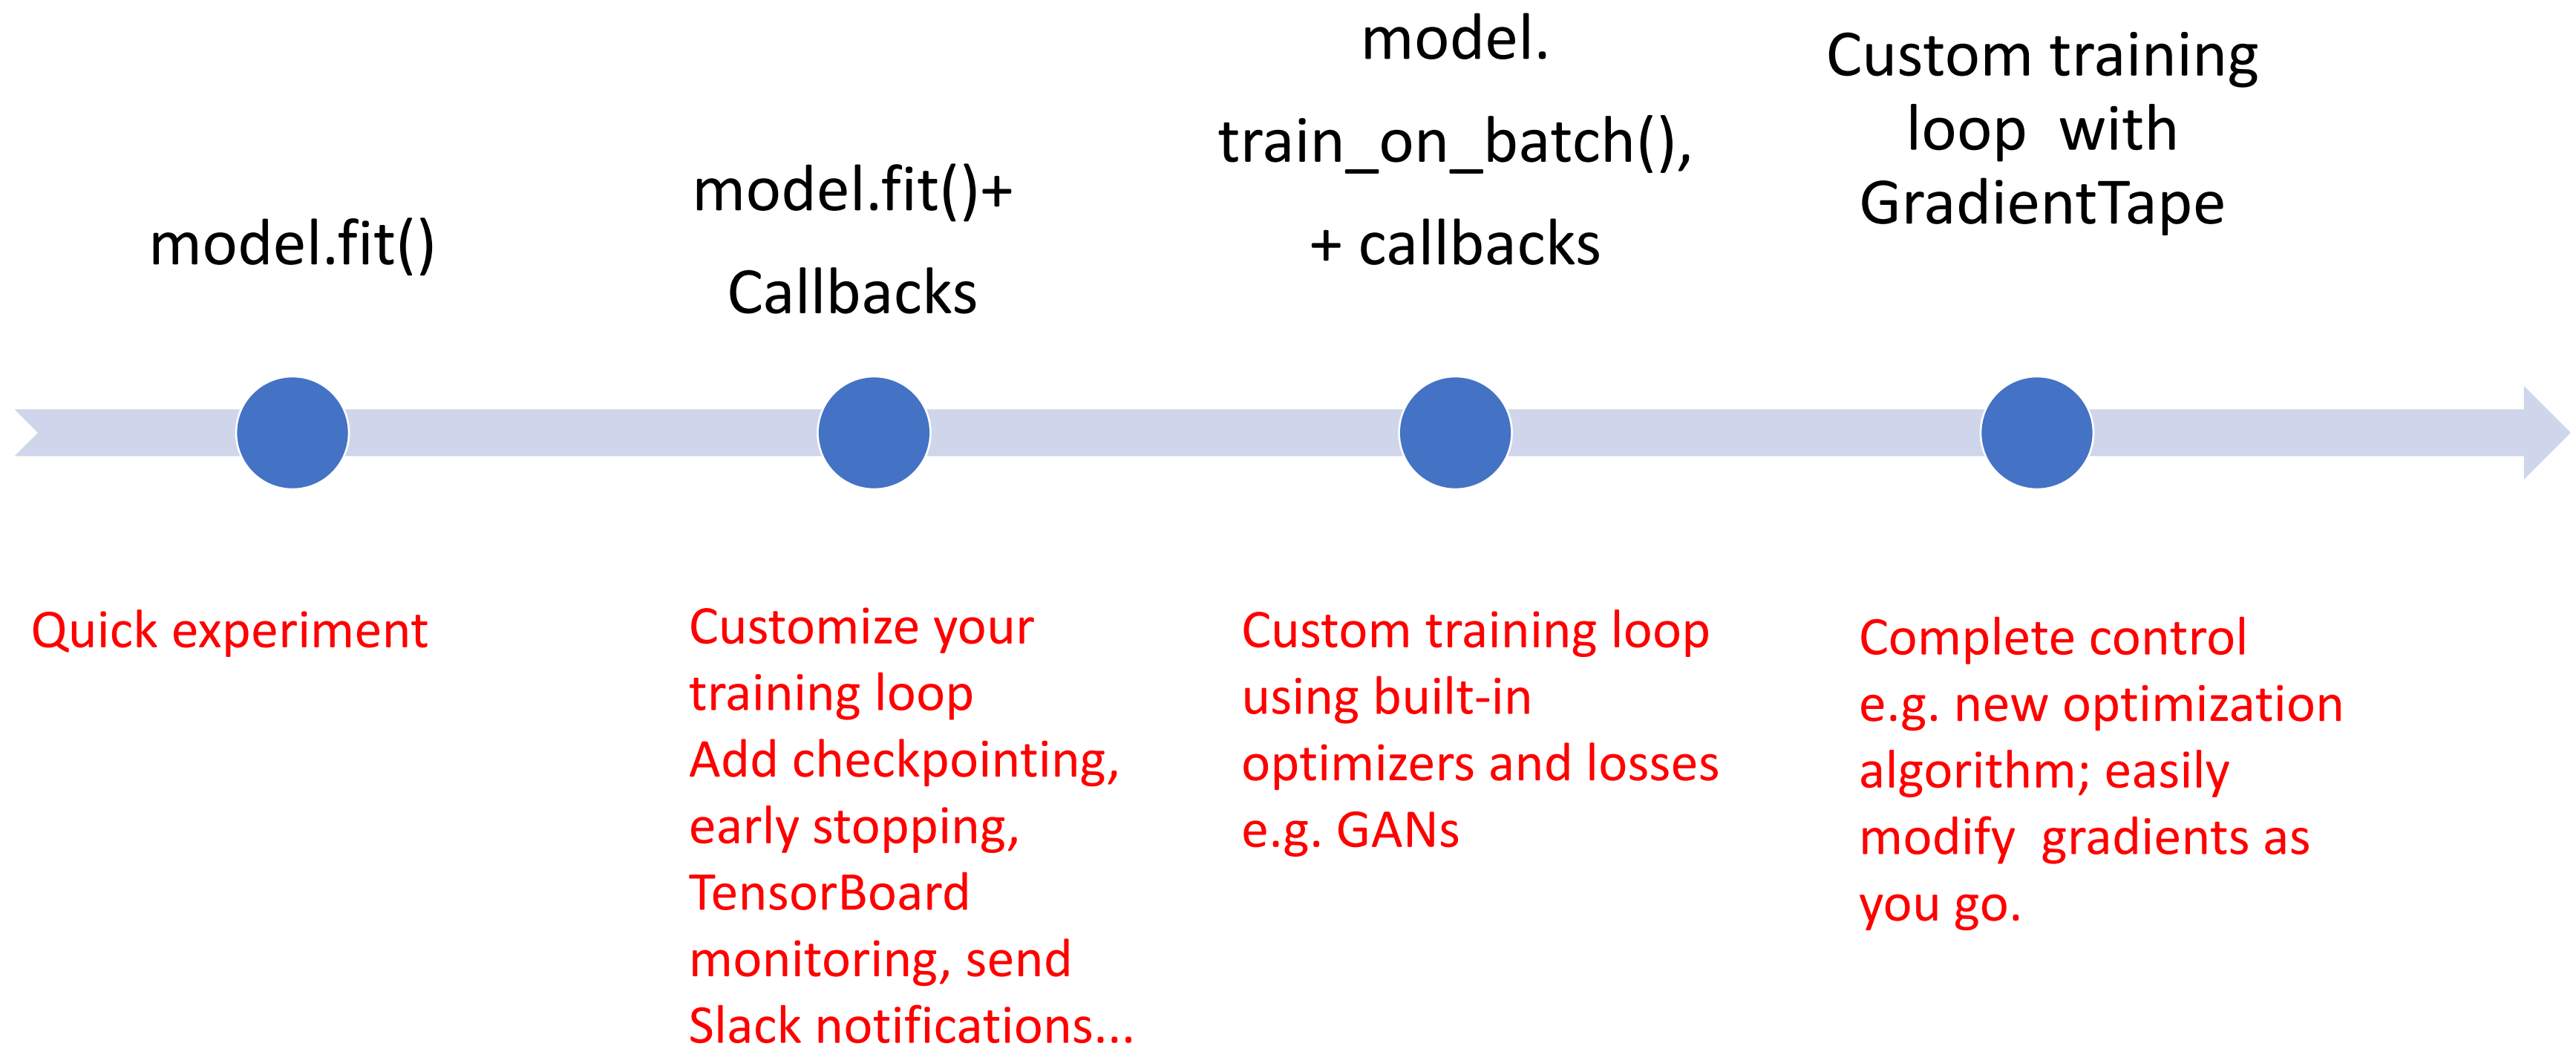


Symbolic vs Imperative APIs

- Symbolic (Keras Sequential)
 - Your model is a graph of layers
 - Any graph you compile will run
 - TensorFlow helps you debug by catching errors at compile time
- Imperative (Keras Subclassing)
 - Your model is Python bytecode
 - Complete flexibility and control
 - Harder to debug / harder to maintain



Model Training





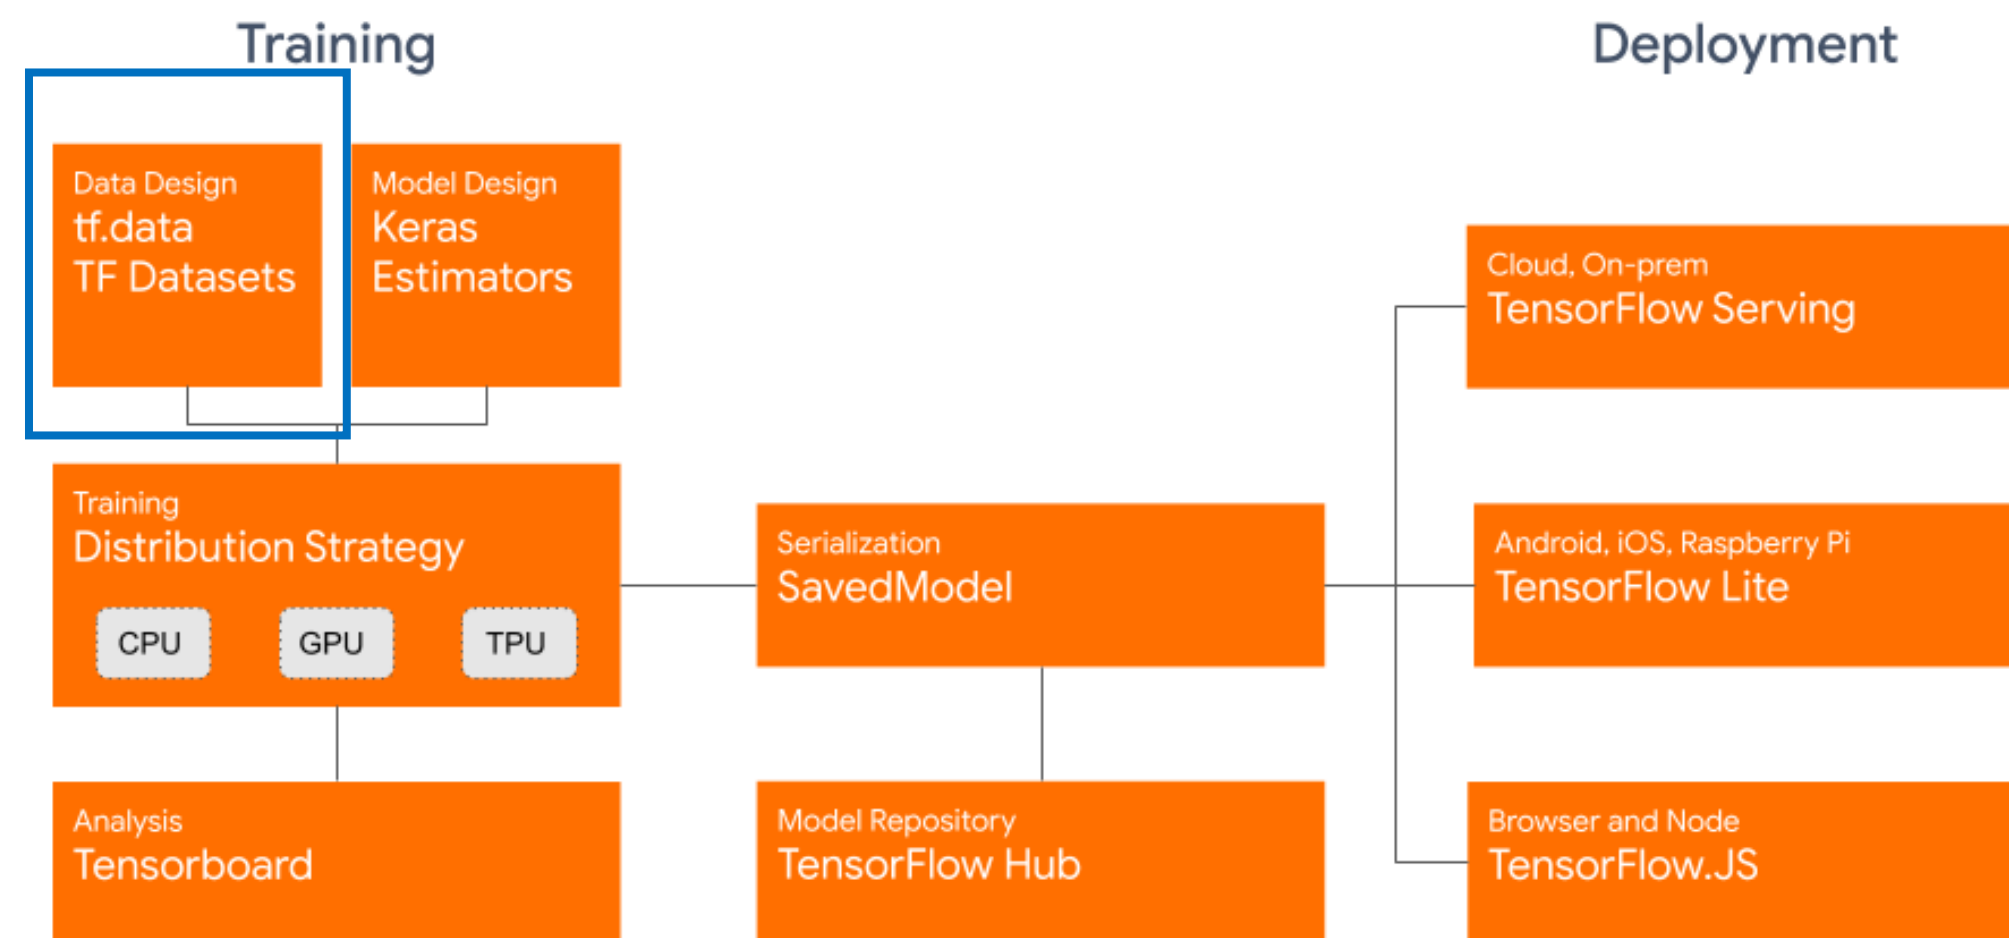
Tensor datatype

- A Tensor has a name, type a rank and a shape
 - The name identifies uniquely the object in the computational graph
 - The type specify the data type, for example `tf.float32` or `tf.int8`.
 - The rank is simply the number of dimensions of the tensor: a scalar has rank 0, a vector has rank 1 and so on.
 - The shape is the number of elements in each dimensions: a scalar has a shape of `()`, a vector a shape of `(d0)`, a matrix shape of `(d0,d1)` and so on (with `d0` and `d1` positive integers).
 - NOTE: the dimension `None` is allowed and indicates an unknown dimension.



Main tensor types

- **tf.Variable*** - will change during training; a variable maintains state in the graph across calls to run()
- **tf.constant**** - will remain constant





```
# Keras datasets
from tensorflow.keras import datasets (train_images,
train_labels), \
    (test_images, test_labels) = datasets.cifar10.load_data()

# TensorFlow Datasets
import tensorflow_datasets as tfds
dataset, metadata = tfds.load('cycle_gan/horse2zebra',
                              with_info=True, as_supervised=True)
```



```
# If you're using TensorFlow Datasets
# Either load your dataset into memory, or
# write a performant input pipeline to load it off disk.
dataset, metadata = tfds.load('mnist',
                              with_info=True, as_supervised=True,
                              in_memory=True)
```



```
# Caching is important to avoid repeated work
# Use either an in-memory cache, or a cache file
def preprocess(img):

    img = tf.cast(image, tf.float32)  img = (img / 127.5) - 1
    img = tf.image.resize(img, [286, 286]) # ...
    return img

image_ds = image_ds.map(
    preprocess, num_parallel_calls=AUTOTUNE).cache()
```

Note: order is important. Cache before shuffling and batching.

Helpful reference (on tf.data, loading images, and caching): [tensorflow.org/tutorials/load_data/images](https://www.tensorflow.org/tutorials/load_data/images)
List of TensorFlow Datasets: [tensorflow.org/datasets/catalog/overview](https://www.tensorflow.org/datasets/catalog/overview)



DEMO

[demo1] <https://colab.research.google.com/drive/1BbMLpUS5-9vnee3DEBq4DKMD2h-cx3cc#scrollTo=4N7XbNDVY8P3>

[demo 2] https://colab.research.google.com/drive/1U1R4fntlQzN93e0WSANgwGtHczwXIbR_#scrollTo=-HJV4JF789aC

[demo 3] https://colab.research.google.com/drive/11AnQ39sHsuUkEC7Lg5__-IT2SLrIEIA8#scrollTo=Y04m-jvKRDsJ



Thank You