

IIT Kharagpur



IIT Madras



IIT Goa



IIT PALAKKAD

IIT Palakkad

APPLIED ACCELERATED ARTIFICIAL INTELLIGENCE

Accelerated TensorFlow

Dr. Satyajit Das

Assistant Professor

Data Science

Computer Science and Engineering

IIT Palakkad



National
Supercomputing
Mission



Centre for
Development of
Advanced Computing




Topics

- 3 ways to create a model
- Accelerated data pipelining
- Distributed training




3 ways to create model

1. Sequential API



2. Functional API



3. Model Subclassing

```
tensorflow.keras.Model  
  
class MySimpleNN(Model):  
    ...
```



Quick overview

```
class WideAndDeepModel(keras.Model):
    def __init__(self, units=30, activation="relu", **kwargs):
        super().__init__(**kwargs) # handles standard arguments ( name, etc.)
        self.hidden1 = keras.layers.Dense(units, activation=activation)
        self.hidden2 = keras.layers.Dense(units, activation=activation)
        self.main_output = keras.layers.Dense(1)
        self.aux_output = keras.layers.Dense(1)

    def call(self, inputs):
        input_A, input_B = inputs
        hidden1 = self.hidden1(input_B)
        hidden2 = self.hidden2(hidden1)
        concat = keras.layers.Concatenate([input_A, hidden2])
        main_output = self.main_output(concat)
        aux_output = self.aux_output(hidden2)
        return main_output, aux_output

model = WideAndDeepModel()
```

```
model = Sequential()
```

```
model.add(Dense(4,activation='relu')) ##<----- You don't have to specify
input size. Just define the hidden layers
model.add(Dense(4,activation='relu'))
model.add(Dense(1))## defining the optimiser and loss function
```

```
model.compile(optimizer='adam',loss='mse')## training the model
model.fit(x=X_train,y=y_train, validation_data=(X_test,y_test),
batch_size=128,epochs=400)
```

```
## Creating the layers
```

```
input_layer = Input(shape=(3,))
Layer_1 = Dense(4, activation="relu")(input_layer)
Layer_2 = Dense(4, activation="relu")(Layer_1)
output_layer= Dense(1, activation="linear")(Layer_2)
```

```
##Defining the model by specifying the input and output layers
```

```
model = Model(inputs=input_layer, outputs=output_layer) ## defining the
optimiser and loss function model.compile(optimizer='adam', loss='mse')
```

```
## training the model model.fit(X_train, y_train,epochs=400,
batch_size=128,validation_data=(X_test,y_test))
```



Why input pipeline is important

- data might not fit into memory
- data might require (randomized) pre-processing
- efficiently utilize hardware
- decouple loading + pre-processing from distribution




- Extract:
 - read data from memory / storage
 - parse file format
- Transform:
 - text vectorization
 - image transformations
 - video temporal sampling
 - shuffling, batching,...
- Load:
 - transfer data to the accelerator



The naive approach


```
def benchmark(dataset, num_epochs=2):  
    start_time = time.perf_counter()  
    for epoch_num in range(num_epochs):  
        for sample in dataset:  
            # Performing a training step  
            time.sleep(0.01)  
    print("Execution time:", time.perf_counter() - start_time)
```

```
benchmark(ArtificialDataset())
```




Prefetching

```
benchmark(  
    ArtificialDataset()  
    .prefetch(tf.data.AUTOTUNE)  
)
```



Interleaving

```
benchmark(  
  tf.data.Dataset.range(2)  
  .interleave(  
    lambda _: ArtificialDataset(),  
    num_parallel_calls=tf.data.AUTOTUNE  
  )  
)
```







Sequential vs parallel mapping

```
benchmark(  
  ArtificialDataset()  
  .map(mapped_function)  
)
```

```
benchmark(  
  ArtificialDataset()  
  .map(  
    mapped_function,  
    num_parallel_calls=tf.data.AUTOTUNE  
  )  
)
```






Caching

```
benchmark(  
  ArtificialDataset()  
  .map( # Apply time consuming operations before cache  
        mapped_function  
  ).cache(  
  ),  
  5  
)
```

Cached dataset

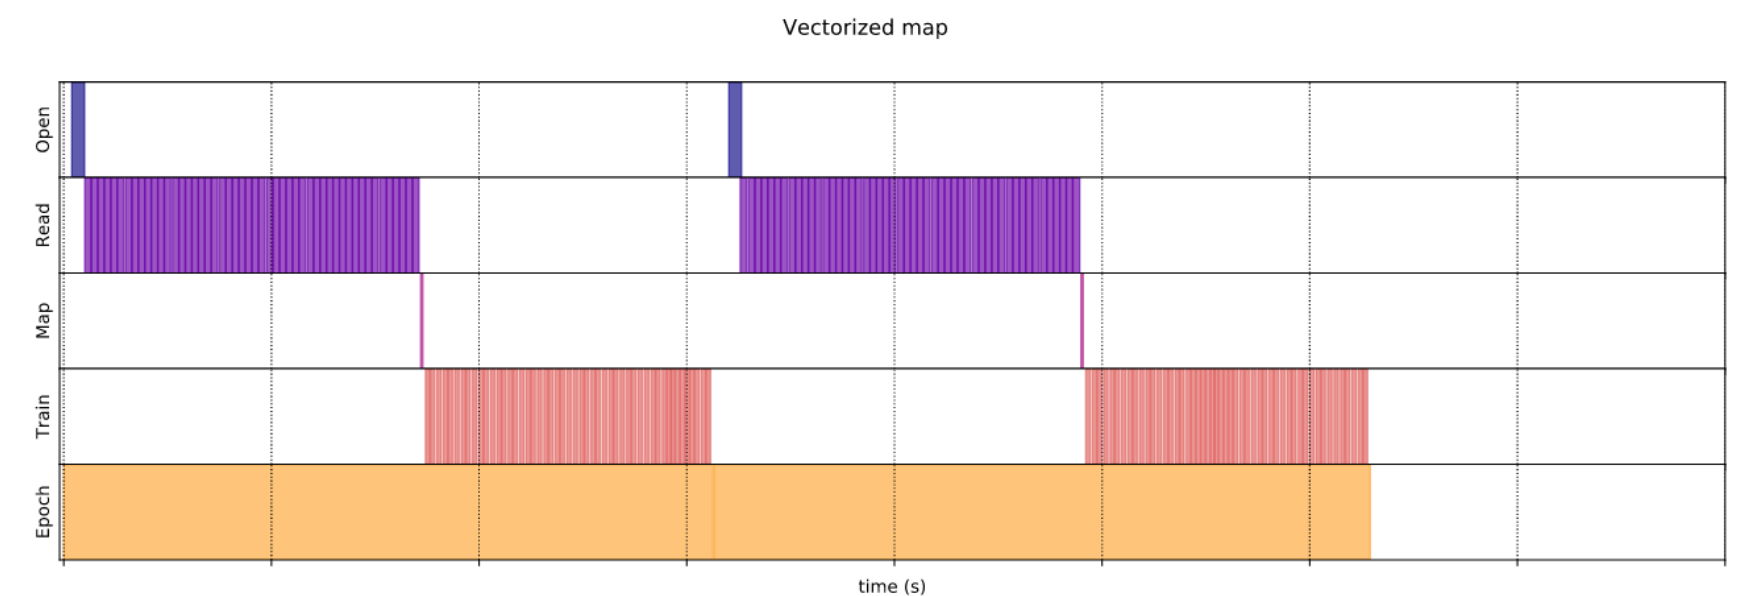
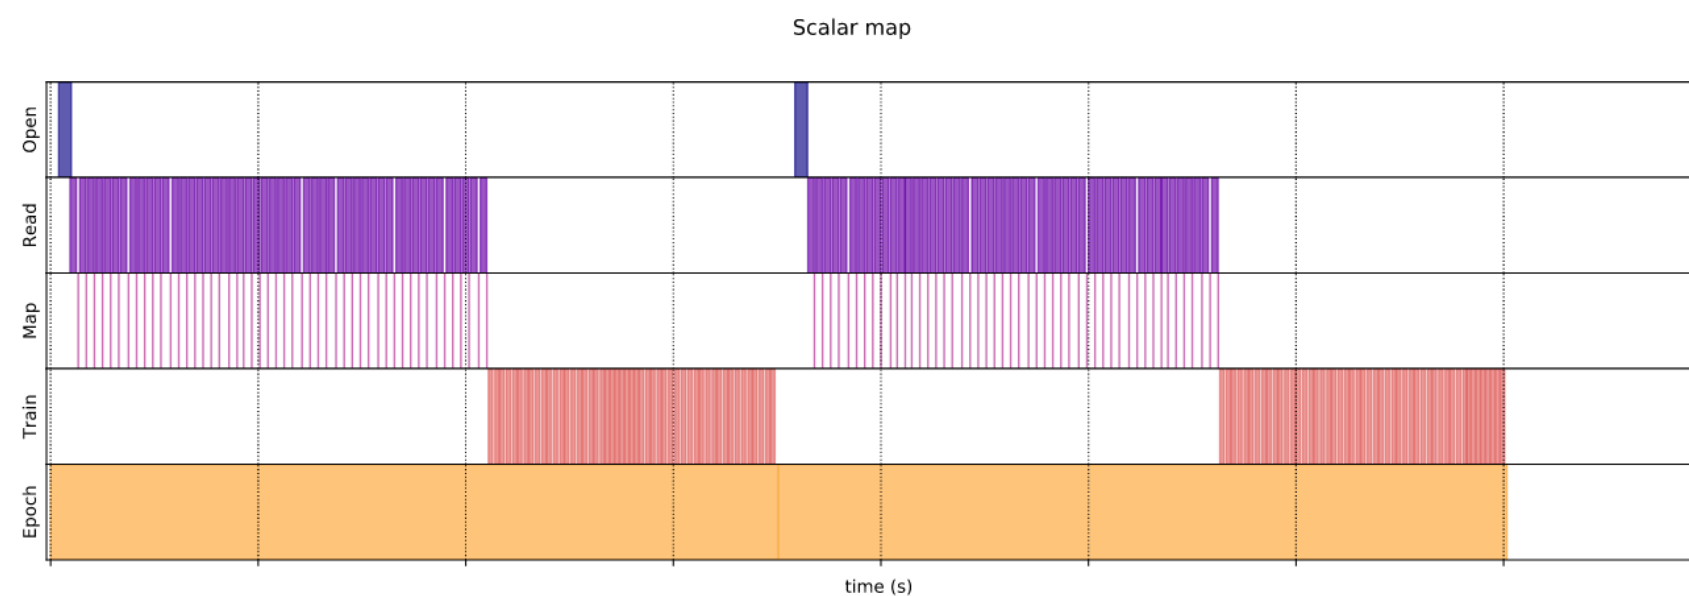




Scalar vs Vectorized mapping

```
fast_benchmark(  
    fast_dataset  
    # Apply function one item at a time  
    .map(increment)  
    # Batch  
    .batch(256)  
)
```

```
fast_benchmark(  
    fast_dataset  
    .batch(256)  
    # Apply function on a batch of items  
    # The tf.Tensor.__add__ method already handle batches  
    .map(increment)  
)
```





TfDS

```
import tensorflow_datasets as tfds

# Download the dataset and create a tf.data.Dataset
ds, info = tfds.load("mnist", split="train", with_info=True)

# Access relevant metadata with DatasetInfo
print(info.splits["train"].num_examples)
print(info.features["label"].num_classes)

# Build your input pipeline
ds = ds.batch(128).repeat(10)

# And get NumPy arrays if you'd like
for ex in tfds.as_numpy(ds):
    np_image, np_label = ex["image"], ex["label"]
```



Why Distributed Training

- An easy way to distribute TensorFlow training.
- Goals:
 - Easy to use - minimal code changes
 - Great out-of-the-box performance
 - Versatile - works with different architectures, hardware and APIs




Use cases of `tf.distribute.strategy`

- Distribute my model using Keras / Estimator API
- Distribute my model using a custom training loop
- Make my layer / library / infrastructure distribute-aware
- Make a new strategy

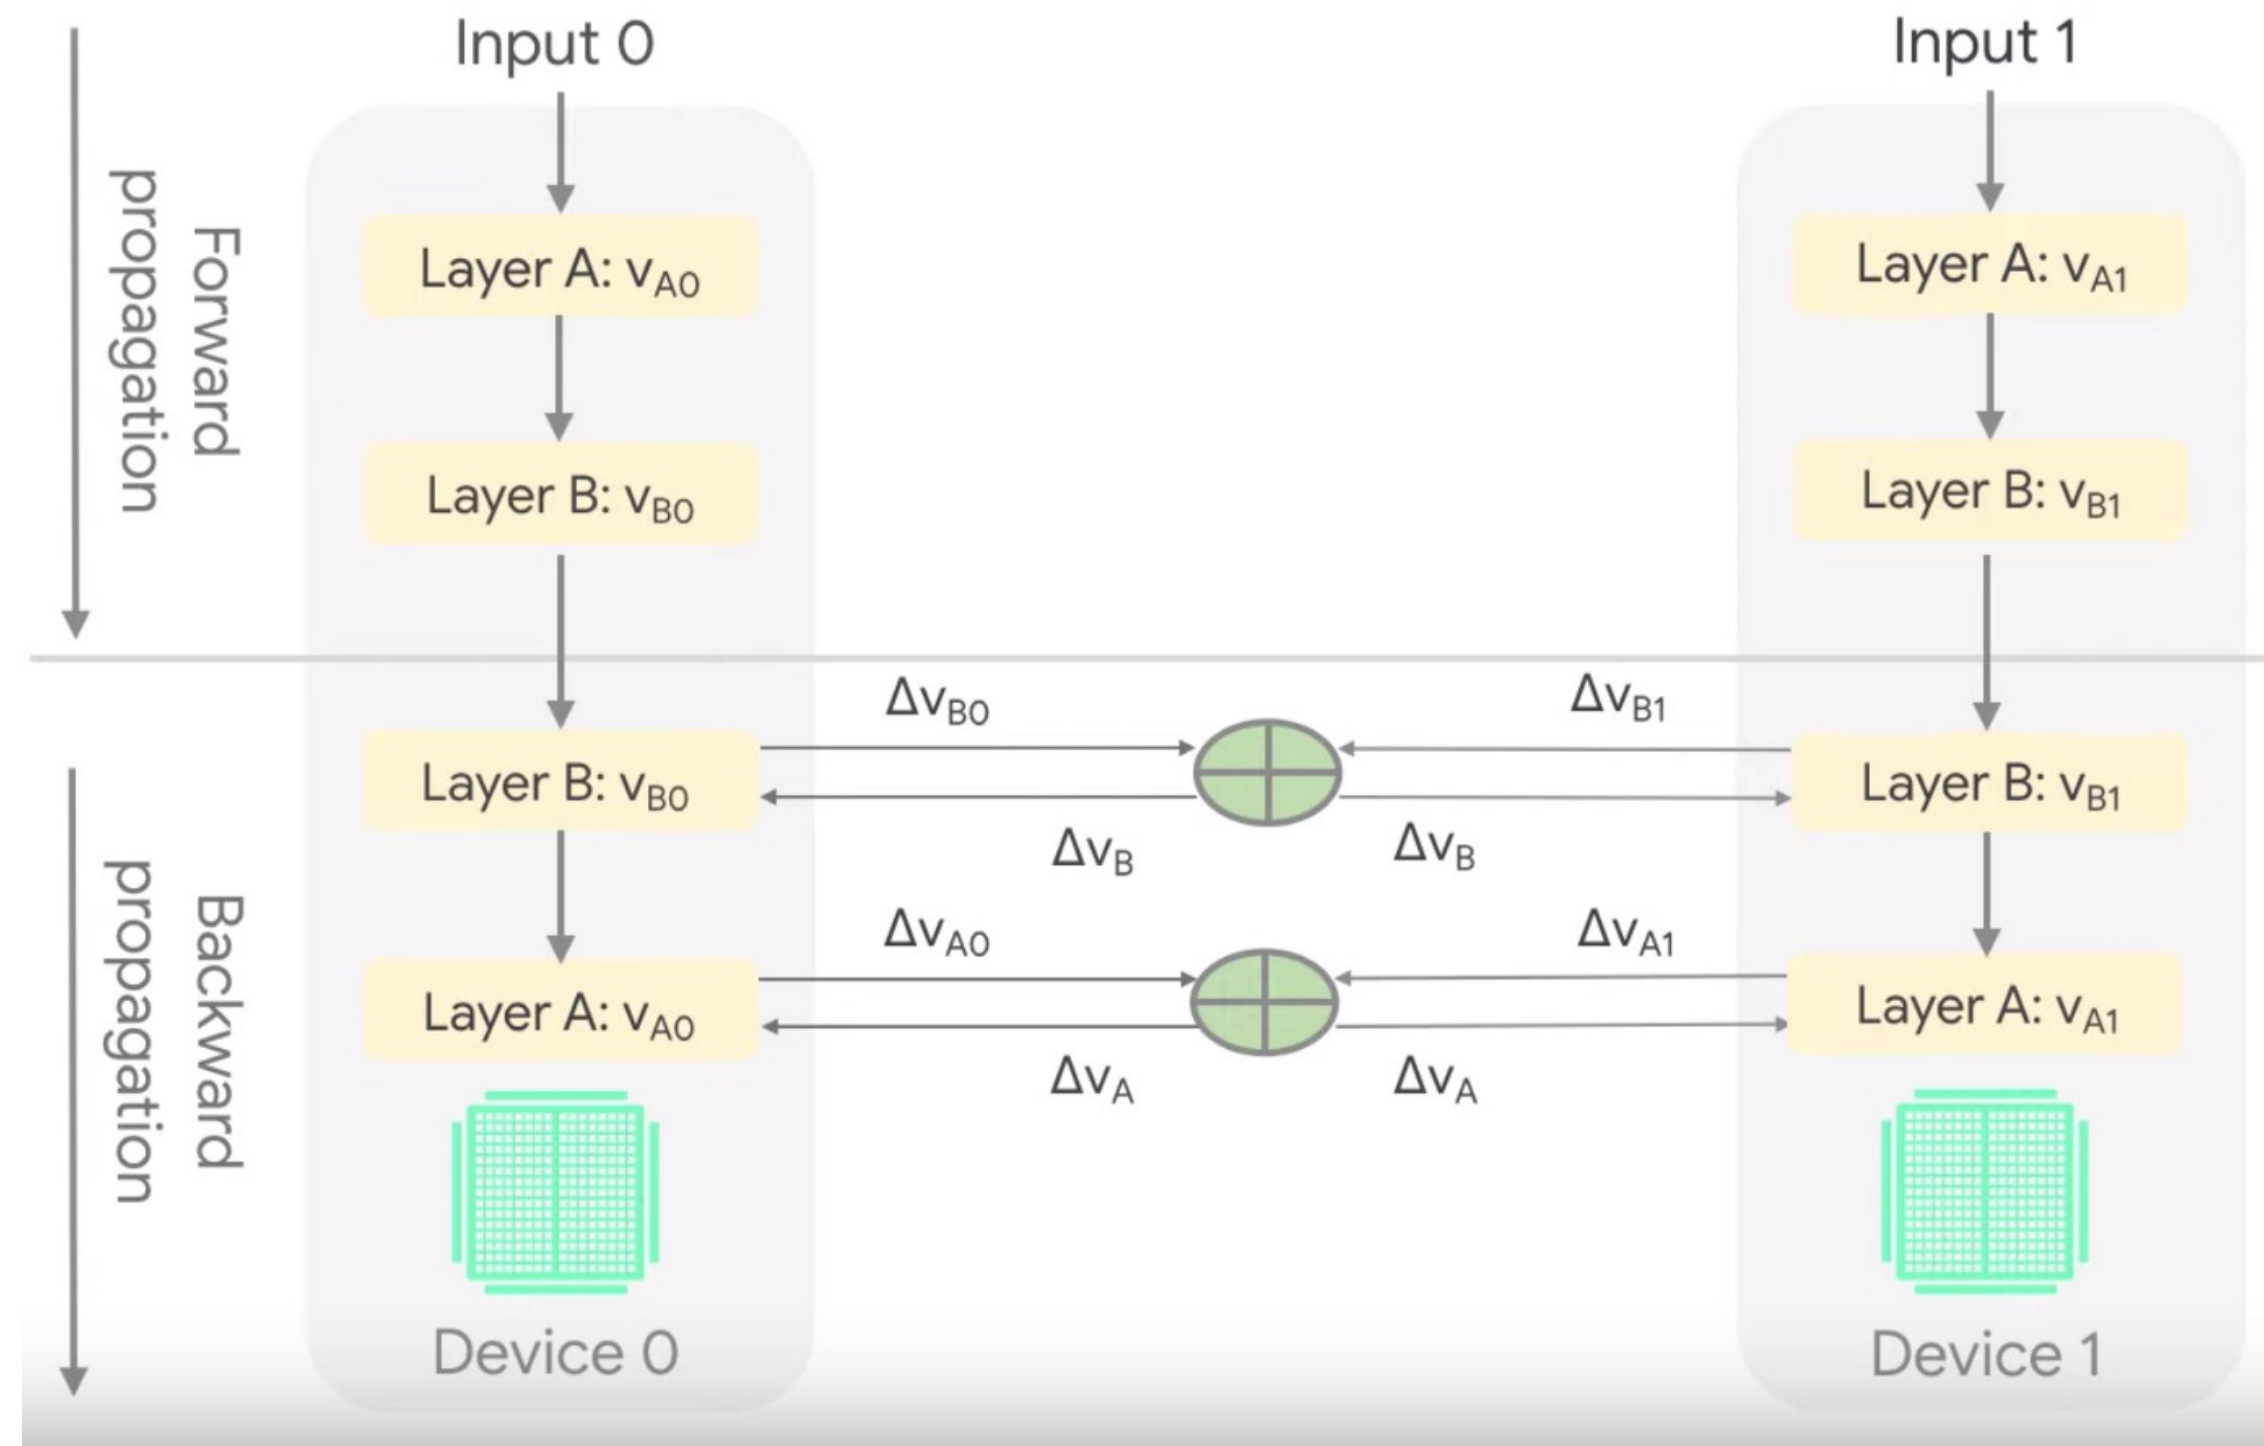
Different API surfaces for each

Multi-GPU All-reduce sync training

- All-reduce synchronous training for multi-GPU
 - replicas are run in lock-step, synchronizing gradients at each step
 - Variables are mirrored on each GPU
 - All-reduce: network efficient way to aggregate gradients



Synchronous training





MirroredStrategy

```
import tensorflow as tf

strategy = tf.distribute.MirroredStrategy()
strategy = tf.distribute.MirroredStrategy(devices=["gpu:0", "gpu:1"])
strategy = tf.distribute.MirroredStrategy(
    cross_device_op=tf.distribute.NcclAllReduce(num_packs=2))
```




Use the strategy with Keras API

```
model = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, 3, activation='relu', input_shape=(28, 28, 1)),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(10)
])

model.compile(loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              optimizer=tf.keras.optimizers.Adam(),
              metrics=['accuracy'])
```

Multi worker all-reduce sync training

- Uses new collective ops
- Workers are run in lock-step,
- synchronizing gradients at each step






Multi worker all-reduce sync training

```
import tensorflow as tf

strategy = tf.distribute.experimental.MultiWorkerMirroredStrategy()
strategy = tf.distribute.experimental.MultiWorkerMirroredStrategy(
    tf.distribute.experimental.CollectiveCommunication.NCCL)
```


All-reduce sync training for TPUs

- Similar to MirroredStrategy
- Uses cross-replica-sum on TPUs to do all reduce



https://medium.com/@decoded_cipher/tensor-processing-units-both-history-and-applications-b3479d92a61d


```
resolver = tf.distribute.cluster_resolver.TPUClusterResolver(tpu='')  
tf.config.experimental_connect_to_cluster(resolver)  
tf.tpu.experimental.initialize_tpu_system(resolver)  
strategy = tf.distribute.experimental.TPUStrategy(resolver)
```



Parameter Servers and Workers


```
ps_strategy = tf.distribute.experimental.ParameterServerStrategy()  
parameter_server_strategy = tf.distribute.experimental.ParameterServerStrategy()
```

Each Averages Portion of the Gradients



Central storage

- 1-machine multi-GPU
 - one copy of each variable on CPU
 - one replica of the model per GPU
 - Can handle embeddings that won't fit on a GPU



```
import tensorflow as tf

strategy = tf.distribute.experimental.CentralStorageStrategy()
```




Reference code

- https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/tutorials/distribute/custom_training.ipynb
- https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/tutorials/distribute/multi_worker_with_keras.ipynb
- https://github.com/python-engineer/tensorflow-course/blob/master/07_Functional_API_Project.ipynb

Thank You